

07- Transparencies in a DDBMS

**School of Computer Science
University of Windsor**

Dr Shafaq Khan

Agenda

- Lecture
 - Various type of Transparencies
- Milestone presentations

Announcements

- Next week – **Lab 3 (graded)**
- **Test 1 – Saturday, July 8th (9 am to 10 am); Location: ER1120**



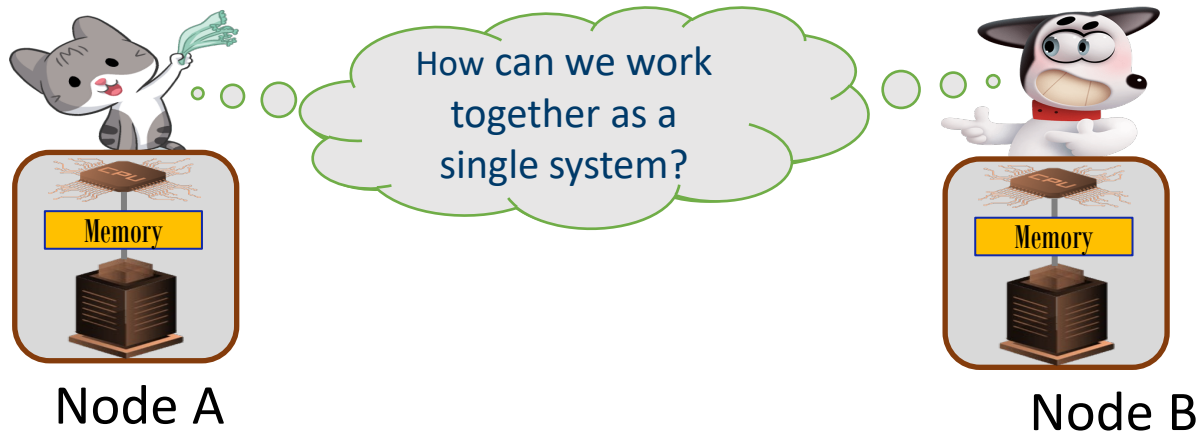
Introductory Questions

What do you mean by transparencies in DDBMS?

What are the different type of transparencies in DDBMS?

Introduction

- From the definition of the DDBMS, the system is expected to make the distribution **transparent** (invisible) to the user.
 - A distributed database is split into fragments that can be stored on different computers and perhaps replicated, should be hidden from the user.
 - The objective of transparency is to make the distributed system appear like a centralized system. This is sometimes referred to as the fundamental principle of distributed DBMSs.



Main types of transparency in a DDBMS



1. Distribution

Fragmentation

Location

Replication

Local Mapping

Naming



2. Transaction

Concurrency

Failure



3. Performance



4. DBMS



1. Distribution Transparency

Distribution transparency allows the user to perceive the database as a single, logical entity. If a DDBMS exhibits distribution transparency, then the user does not need to know the data is fragmented (fragmentation transparency) or the location of data items (location transparency).



1.1 Fragmentation Transparency

The end user or programmer does not need to know that a database is partitioned.



1.2 Location Transparency

The end user or programmer does not need to know the location of data items.



1.3 Replication Transparency

The user is unaware of the replication of fragments.



1.4 Local Mapping Transparency

The user needs to specify both fragment names and the location of data items, taking into consideration any replication that may exist.



1.5 Naming Transparency

The DDBMS must ensure that no two sites create a database object with the same name.



Distribution Transparency : Example

Consider the following example to illustrate we consider the distribution of the Staff such that:

$$S_1 = \Pi_{\text{staffNo, position, sex, DOB, salary}}(\text{Staff}) \quad \text{located at site 5}$$
$$S_2 = \Pi_{\text{staffNo, fName, lName, branchNo}}(\text{Staff})$$

$$S_{21} = \sigma_{\text{branchNo}='B003'}(S_2) \quad \text{located at site 3}$$
$$S_{22} = \sigma_{\text{branchNo}='B005'}(S_2) \quad \text{located at site 5}$$
$$S_{23} = \sigma_{\text{branchNo}='B007'}(S_2) \quad \text{located at site 7}$$

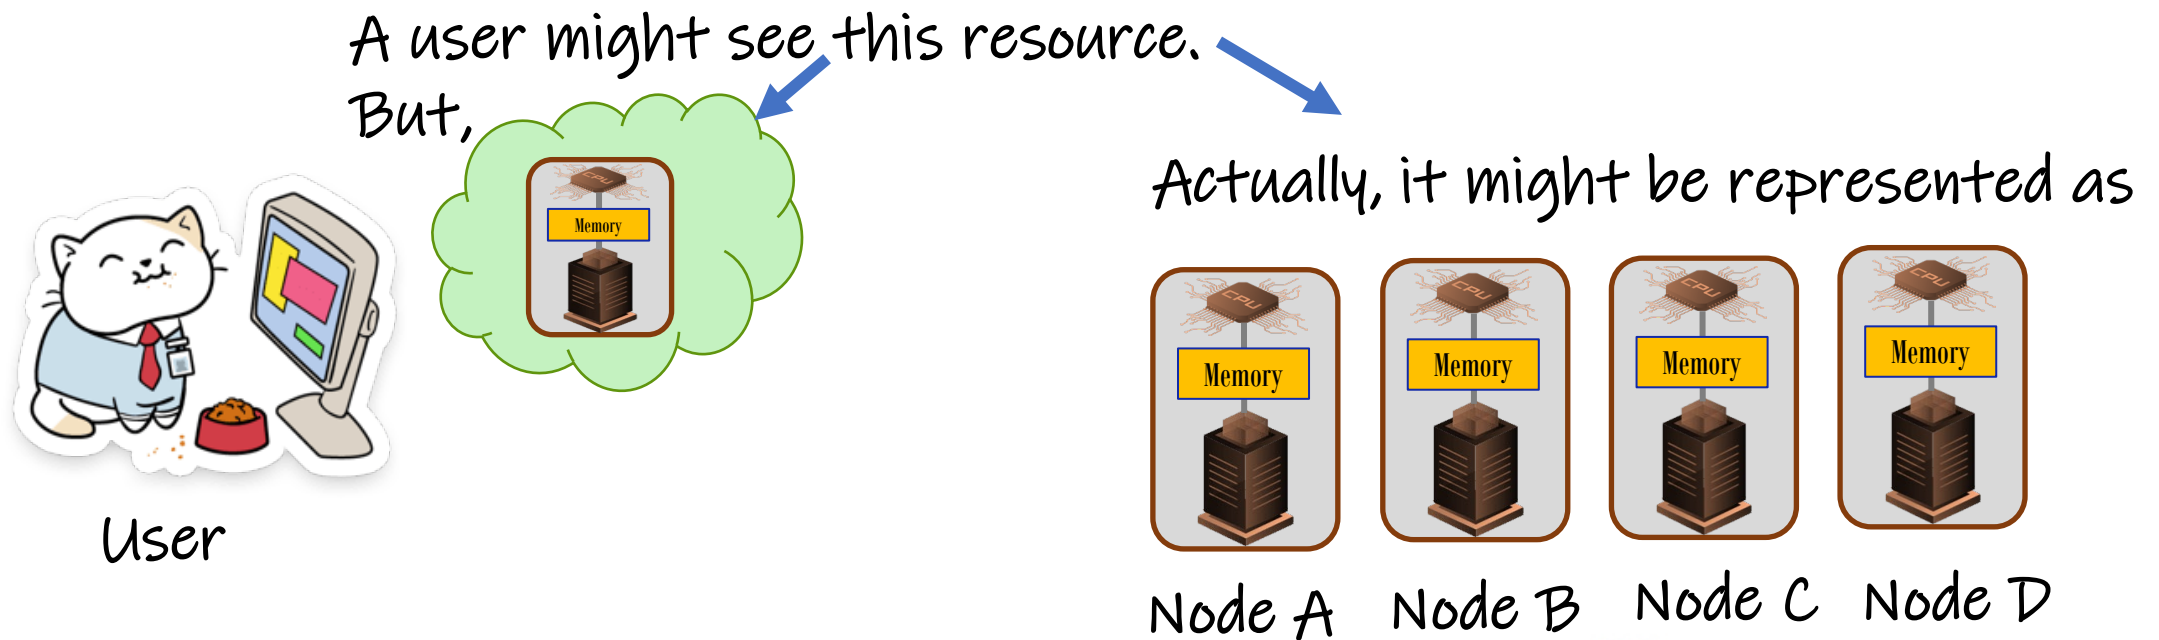


1.1 Fragmentation transparency

Fragmentation is the **highest level** of distribution transparency.

If fragmentation transparency is provided by the DDBMS, then the user does not need to know that the data is fragmented.

As a result, database accesses are based on the global schema, so the user does not need to specify **fragment names or data locations**.



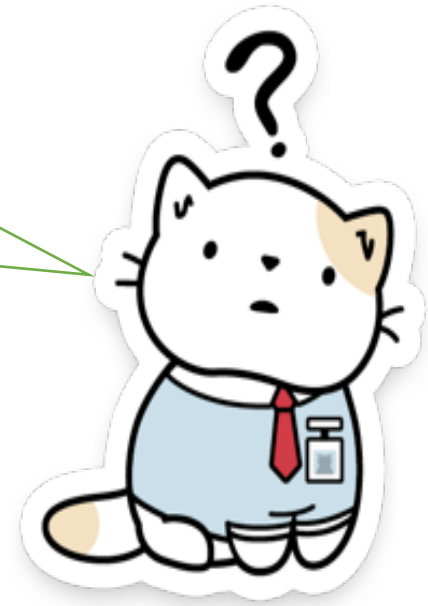
1.1 Fragmentation transparency (cont.)

Example: Retrieve the names of all managers

Note: with fragmentation transparency we could write:

```
SELECT fName, lName  
FROM Staff  
WHERE position = 'Manager';
```

How will you write
the same query in a
centralized system?



1.2 Location transparency

Location is the **middle level** of distribution transparency.

With location transparency, the user **must know** how the data has been fragmented but still **does not have to know** the location of the data.

The main advantage of location transparency is that the database may be physically reorganized without impacting on the application programs that access them.

Example: Retrieve the names of all managers

```
SELECT fName, lName FROM S21
WHERE staffNo IN (SELECT staffNo FROM S1 WHERE position = 'Manager')
UNION
SELECT fName, lName FROM S22
WHERE staffNo IN (SELECT staffNo FROM S1 WHERE position = 'Manager')
UNION
SELECT fName, lName FROM S23
WHERE staffNo IN (SELECT staffNo FROM S1 WHERE position = 'Manager');
```



1.3 Local mapping transparency

The **lowest level** of distribution transparency.

The user needs to specify both **fragment names** and the **location** of data items, taking into consideration any replication that may exist.

```
SELECT fName, lName FROM S21 AT SITE 3  
WHERE staffNo IN (SELECT staffNo FROM S1 AT SITE 5 WHERE position = 'Manager')  
UNION  
SELECT fName, lName FROM S22 AT SITE 5  
WHERE staffNo IN (SELECT staffNo FROM S1 AT SITE 5 WHERE position = 'Manager')  
UNION  
SELECT fName, lName FROM S23 AT SITE 7  
WHERE staffNo IN (SELECT staffNo FROM S1 AT SITE 5 WHERE position = 'Manager');
```

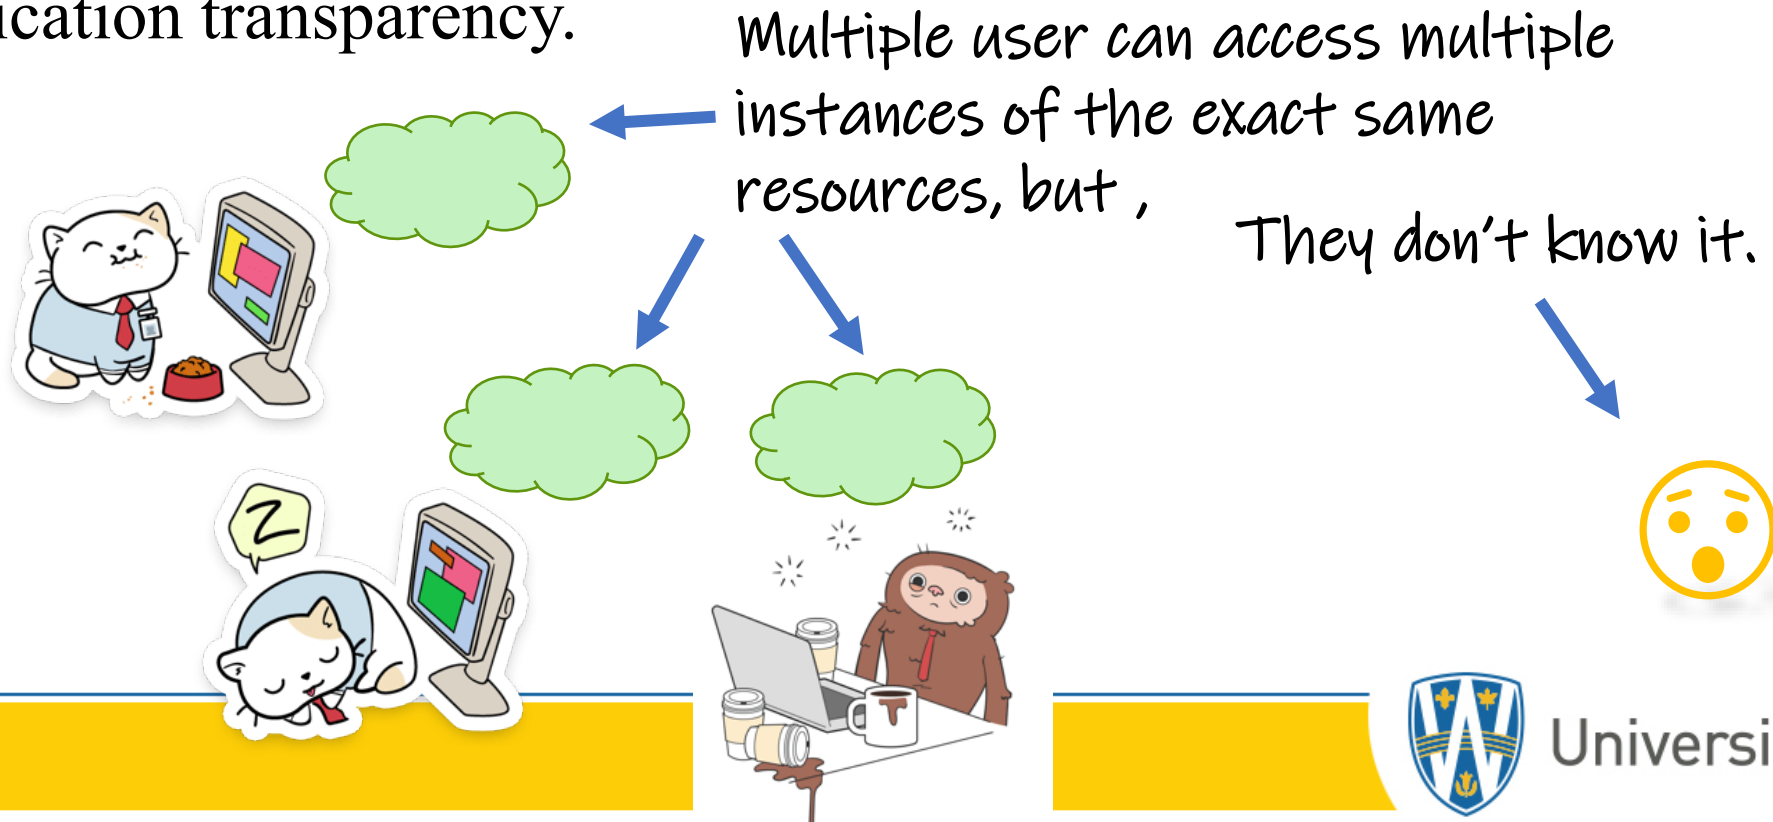


1.4 Replication transparency

The user is unaware of the replication of fragments.

Replication transparency is implied by location transparency.

However, it is possible for a system not to have location transparency but to have replication transparency.



1.5 Naming transparency (1/3)

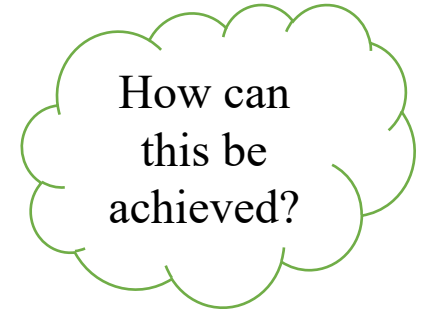
Each item in a DDB must have a unique name.

DDBMS must ensure that no two sites create a database object with same name.

Solution: create a **central name server**.

However, this results in:

- loss of some local autonomy;
- central site may become a bottleneck;
- low availability; if the central site fails, remaining sites cannot create any new objects.



1.5 Naming transparency (2/3)

Alternative solution - prefix object with identifier of site that created it.
For example,

Branch created at site S1 might be named **S1.BRANCH**.

Also need to identify each fragment and its copies.

Thus, copy 2 of fragment 3 of Branch created at site S1 might be referred to as **S1.BRANCH.F3.C2**.

However, this results in loss of distribution transparency.

To resolve the problems with both these solutions, use **aliases** for each database object.

For example,

S1.Branch.F3.C2 might be known as **LocalBranch** by the user at site S1.

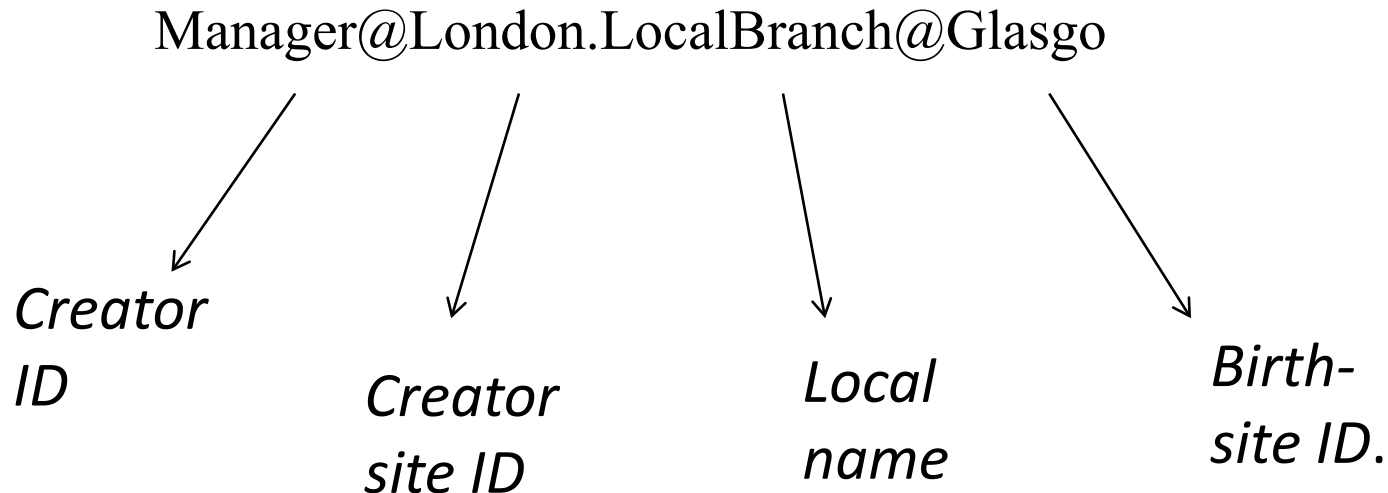
DDBMS has task of mapping an alias to appropriate database object.



1.5 Naming transparency (3/3)

The distributed system **R*** distinguishes between an object's **printname** and its **systemwide name**.

- ✓ The **printname** is the name that the **users** normally use to refer to the object.
- ✓ The **systemwide** name is a globally unique **internal identifier** for the object that is guaranteed never to change.
- ✓ For example, the systemwide name:



1.5 Naming transparency (3/3)

Systemwide Name

- *Creator ID*: A unique site identifier for the user who created the object;
 - *Creator site ID*: A globally unique identifier for the site from which the object was created
 - *Local name*: An unqualified name for the object;
 - *Birth-site ID*: A globally unique identifier for the site at which the object was initially stored
-
- DDBMS Data Object Name:
 - *Creator ID@ Creator site ID. Local name@ Birth-site ID*
 - Fragment S1 stored on Site 5 : [Manager@Site5.S1@Site5](#)
 - Fragment S21 stored on Site 3: [Manager@Site5.S21@Site3](#)
 - Fragment S22 stored on Site 5: [Manager@Site5.S22@Site5](#)
 - Fragment S23 stored on Site 7: [Manager@Site5.S23@Site7](#)
 - Replica of fragment S23 on Site 9: [Manager@Site7.S23@Site9](#)
 - The Global Schema Catalog (GSC) stores all the naming information



2. Transaction Transparency

- A site may have a number of **local** and **global transactions** that are being executed concurrently
- Local transactions have to **synchronize** with other local transactions
- Local transactions have to synchronize with global transactions both on the site and on other sites



2. Transaction Transparency: Example

Consider a transaction T that prints out names of all staff, using schema defined earlier as S_1 , S_2 , S_{21} , S_{22} , and S_{23} . Define three sub-transactions T_{S_3} , T_{S_5} , and T_{S_7} to represent agents at sites 3, 5, and 7.

Time	T_{S_3}	T_{S_5}	T_{S_7}
t_1	begin_transaction	begin_transaction	begin_transaction
t_2	read(fName, lName)	read(fName, lName)	read(fName, lName)
t_3	print(fName, lName)	print(fName, lName)	print(fName, lName)
t_4	end_transaction	end_transaction	end_transaction

Distributed transaction.

DDBMS ensure synchronization of sub-transactions with other local transactions that are executing concurrently at a site, but it must also ensure synchronization of sub-transactions with global transactions running simultaneously at the same or different sites.



Challenges

- Transaction Transparency is complicated by
 - Fragmentation
 - Replication
 - Allocation schemas

2. Transaction Transparency..

- Concurrency Transparency
- Failure Transparency



Concurrency Transparency

DDBMS must ensure

- ✓ both global and local transactions do not interfere with each other.
- ✓ consistency of all sub-transactions of global transaction.

Replication makes concurrency more complex.

- ✓ If a copy of a replicated data item is updated, update must be propagated to all copies.
- ✓ Could propagate changes as part of original transaction, making it an atomic operation.
- ✓ However, if one site holding copy is not reachable, then transaction is delayed until site is reachable.



Failure Transparency

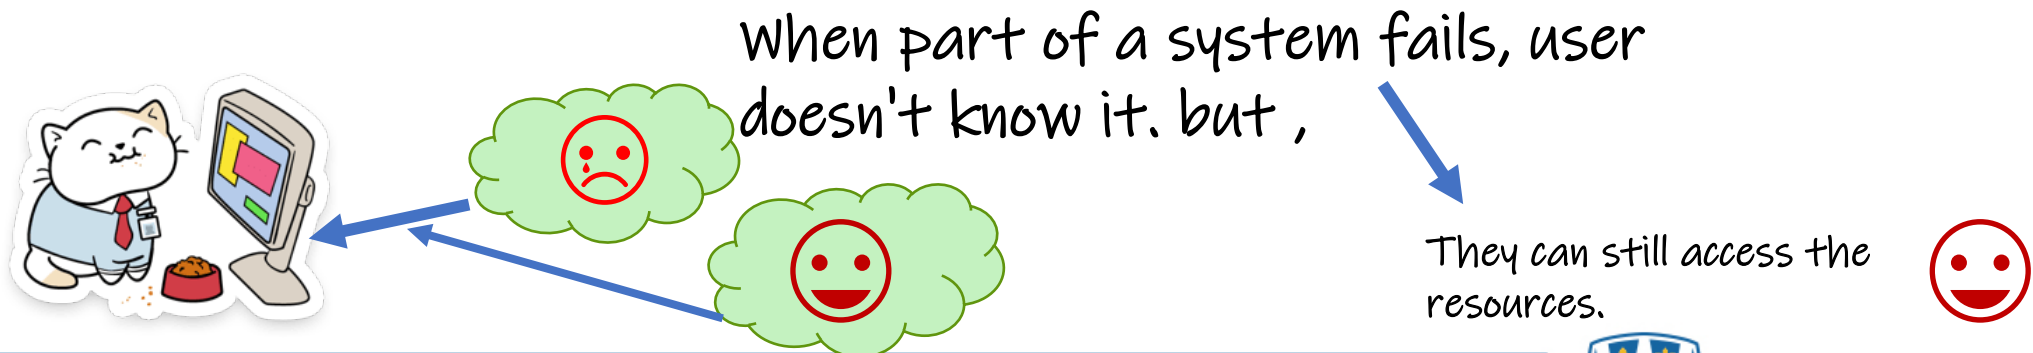
DDBMS must ensure **atomicity** and **durability** of global transaction.

- ✓ Means ensuring that sub-transactions of global transaction either all commit or all abort.

In the distributed environment, the DDBMS must also cater for:

- ✓ the loss of a message;
- ✓ the failure of a communication link;
- ✓ the failure of a site;
- ✓ network partitioning.


Thus, DDBMS must synchronize global transaction to ensure that all sub-transactions have completed successfully before recording a final COMMIT for global transaction.



Example

Consider a global transaction that has to update data at two sites, say S1 and S2.

The subtransaction at site S1 completes successfully and commits



The distributed database is now in an inconsistent state: we are unable to uncommit the data at site S1

The subtransaction at site S2 is unable to commit and rolls back the changes to ensure local consistency



3. Performance Transparency

Optimal database performance with minimal associated costs is the goal of query optimization.

DDBMS must perform as if it were a centralized DBMS.

- ✓ DDBMS should **not suffer any performance degradation** due to distributed architecture.
- ✓ DDBMS should determine **most cost-effective strategy** to execute a request

Distributed Query Processor (DQP) maps data request into ordered sequence of operations on local databases. Must consider fragmentation, replication, and allocation schemas.

DQP has to decide:

- ✓ which fragment to access;
- ✓ which copy of a fragment to use;
- ✓ which location to use.



Performance Transparency

The DQP produces an execution strategy that is optimized with respect to some cost function.

Typically, the costs associated with a distributed request include:

- ✓ the **access time (I/O) cost** involved in accessing the physical data on disk;
- ✓ the **CPU time cost** incurred when performing operations on data in main memory;
- ✓ the **communication cost** associated with the transmission of data across the network.



Performance Transparency – Example

Property(propNo, city)	10000 records in London
Client(clientNo,maxPrice)	100000 records in Glasgow
Viewing(propNo, clientNo)	1000000 records in London

Query: List the properties in Aberdeen that have been viewed by clients who have a maximum price limit greater than £200,000

```
SELECT p.propNo
FROM Property p INNER JOIN (Client c INNER JOIN Viewing v ON c.clientNo = v.clientNo) ON p.propNo = v.propNo
WHERE p.city='Aberdeen' AND c.maxPrice > 200000;
```

Assume:

Each tuple in each relation is 100 characters long.

10 renters with maximum price greater than £200,000.

100 000 viewings for properties in Aberdeen.

Computation time negligible compared to communication time.

communication system has a data transmission rate of 10,000 characters per second and a 1-second access delay to send a message from one site to another.



Performance Transparency – Example

Strategy 1:	Move the Client relation to London and process query there: Time = $1 + (100\,000 * 100/10\,000) \cong 16.7$ minutes
Strategy 2:	Move the Property and Viewing relations to Glasgow and process query there: Time = $2 + [(1\,000\,000 + 10\,000) * 100/10\,000] \cong 28$ hours
Strategy 3:	Join the Property and Viewing relations at London, select tuples for Aberdeen properties and then for each of these tuples in turn check at Glasgow to determine whether the associated client's maxPrice > £200,000. The check for each tuple involves two messages: a query and a response. Time = $100\,000 * (1 + 100/10\,000) + 100\,000 * 1 \cong 2.3$ days
Strategy 4:	Select clients with maxPrice > £200,000 at Glasgow and for each one found, check at London to see whether there is a viewing involving that client and an Aberdeen property. Again, two messages are needed: Time = $10 * (1 + 100/10\,000) + 10 * 1 \cong 20$ seconds
Strategy 5:	Join Property and Viewing relations at London, select Aberdeen properties, project result over propertyNo and clientNo, and move this result to Glasgow for matching with maxPrice > £200,000. For simplicity, we assume that the projected result is still 100 characters long: Time = $1 + (100\,000 * 100/10\,000) \cong 16.7$ minutes
Strategy 6:	Select clients with maxPrice > £200,000 at Glasgow and move the result to London for matching with Aberdeen properties: Time = $1 + (10 * 100/10\,000) \cong 1$ second



4. DBMS Transparency

DBMS transparency hides the knowledge that the local DBMSs may be different and is therefore applicable only to heterogeneous DDBMSs.

It is one of the most difficult transparencies to provide as a generalization.



Summary

Distribution Transparency : Fragmentation, Location, Replication, Local Mapping and Naming

Transaction Transparency: Concurrency Transparency , Failure Transparency

Performance Transparency

DBMS Transparency



Any Questions

